

hand-gesture-recognition

September 4, 2023

0.1 Libraries

```
[16]: import fastai
      from fastai.vision.all import *
      import numpy as np
      import matplotlib.pyplot as plt
      import tensorflow as tf
      import pandas as pd
```

0.2 Load files

```
[17]: PATH = "/kaggle/input/leapgestrecog/leapGestRecog"
      RANDOM_STATE = 42
      TEST_SIZE = 0.3
      IMG_SIZE = 128
```


0.2.1 Create Datablock

```
[18]: data = DataBlock(
      blocks=(ImageBlock, CategoryBlock),
      get_items = get_image_files,
      splitter=RandomSplitter(valid_pct=TEST_SIZE, seed=RANDOM_STATE),
      get_y = parent_label,
      item_tfms = Resize(IMG_SIZE)
      )
```

```
[19]: dataloaders = data.dataloaders(PATH, bs=64)
```

0.2.2 Sample from data files

```
[20]: dataloaders.show_batch()
```



```
[21]: learn = vision_learner(dataloaders, resnet34, metrics = error_rate)
```


```
/opt/conda/lib/python3.10/site-packages/torchvision/models/_utils.py:208:
UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be
removed in the future, please use 'weights' instead.
  warnings.warn(
/opt/conda/lib/python3.10/site-packages/torchvision/models/_utils.py:223:
UserWarning: Arguments other than a weight enum or `None` for 'weights' are
deprecated since 0.13 and may be removed in the future. The current behavior is
equivalent to passing `weights=ResNet34_Weights.IMAGENET1K_V1`. You can also use
`weights=ResNet34_Weights.DEFAULT` to get the most up-to-date weights.
  warnings.warn(msg)
```

```
[22]: learn.lr_find()
```

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

```
[22]: SuggestedLRs(valley=0.0014454397605732083)
```



```
[23]: learn.fine_tune(4, 3e-2)
```

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>


<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

```
[24]: learn.show_results()
```

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>



```
[25]: interp = Interpretation.from_learner(learn)
      interp.plot_top_losses(9)
```

<IPython.core.display.HTML object>


<IPython.core.display.HTML object>

<IPython.core.display.HTML object>


<IPython.core.display.HTML object>

Prediction/Actual/Loss/Probability


03_fist/02_l / 1.26 / 0.70



02_l/02_l / 0.59 / 0.56




01_palm/01_palm / 0.34 / 0.71




07_ok/07_ok / 0.27 / 0.76




01_palm/01_palm / 0.26 / 0.77




02_l/02_l / 0.23 / 0.80




02_l/02_l / 0.17 / 0.84



01_palm/01_palm / 0.13 / 0.88



02_l/02_l / 0.12 / 0.89



```
[26]: learn.summary()
```

```
<IPython.core.display.HTML object>
```

```
<IPython.core.display.HTML object>
```

```
[26]: Sequential (Input shape: 64 x 3 x 128 x 128)
```

```
=====
Layer (type)           Output Shape          Param #   Trainable
=====
                        64 x 64 x 64 x 64
```

Conv2d	9408	True
BatchNorm2d	128	True
ReLU		

64 x 64 x 32 x 32

MaxPool2d		
Conv2d	36864	True
BatchNorm2d	128	True
ReLU		
Conv2d	36864	True
BatchNorm2d	128	True
Conv2d	36864	True
BatchNorm2d	128	True
ReLU		
Conv2d	36864	True
BatchNorm2d	128	True
Conv2d	36864	True
BatchNorm2d	128	True
ReLU		
Conv2d	36864	True
BatchNorm2d	128	True

64 x 128 x 16 x 16

Conv2d	73728	True
BatchNorm2d	256	True
ReLU		
Conv2d	147456	True
BatchNorm2d	256	True
Conv2d	8192	True
BatchNorm2d	256	True
Conv2d	147456	True
BatchNorm2d	256	True
ReLU		
Conv2d	147456	True
BatchNorm2d	256	True
Conv2d	147456	True
BatchNorm2d	256	True
ReLU		
Conv2d	147456	True
BatchNorm2d	256	True
Conv2d	147456	True
BatchNorm2d	256	True
ReLU		
Conv2d	147456	True
BatchNorm2d	256	True

64 x 256 x 8 x 8

Conv2d	294912	True
BatchNorm2d	512	True
ReLU		
Conv2d	589824	True
BatchNorm2d	512	True
Conv2d	32768	True
BatchNorm2d	512	True
Conv2d	589824	True
BatchNorm2d	512	True
ReLU		
Conv2d	589824	True
BatchNorm2d	512	True
Conv2d	589824	True
BatchNorm2d	512	True
ReLU		
Conv2d	589824	True
BatchNorm2d	512	True
Conv2d	589824	True
BatchNorm2d	512	True
ReLU		
Conv2d	589824	True
BatchNorm2d	512	True
Conv2d	589824	True
BatchNorm2d	512	True
ReLU		
Conv2d	589824	True
BatchNorm2d	512	True
Conv2d	589824	True
BatchNorm2d	512	True
ReLU		
Conv2d	589824	True
BatchNorm2d	512	True
Conv2d	589824	True
BatchNorm2d	512	True
ReLU		
Conv2d	589824	True
BatchNorm2d	512	True

64 x 512 x 4 x 4

Conv2d	1179648	True
BatchNorm2d	1024	True
ReLU		
Conv2d	2359296	True
BatchNorm2d	1024	True
Conv2d	131072	True
BatchNorm2d	1024	True
Conv2d	2359296	True
BatchNorm2d	1024	True
ReLU		
Conv2d	2359296	True
BatchNorm2d	1024	True
Conv2d	2359296	True

BatchNorm2d	1024	True
ReLU		
Conv2d	2359296	True
BatchNorm2d	1024	True

64 x 512 x 1 x 1

AdaptiveAvgPool2d
AdaptiveMaxPool2d

64 x 1024

Flatten		
BatchNorm1d	2048	True
Dropout		

64 x 512

Linear	524288	True
ReLU		
BatchNorm1d	1024	True
Dropout		

64 x 10

Linear	5120	True
--------	------	------

Total params: 21,817,152
Total trainable params: 21,817,152
Total non-trainable params: 0

Optimizer used: <function Adam at 0x7a0ed796c3a0>
Loss function: FlattenedLoss of CrossEntropyLoss()

Model unfrozen

Callbacks:

- TrainEvalCallback
- CastToTensor
- Recorder
- ProgressCallback

0.2.3 Exported model

```
[27]: learn.export("model-r34")
```